



Chapter 2. Multi-armed Bandits

Reinforcement Learning

발표자 정의정



목차

Chapter 02

Multi-armed Bandits

2.0 Introduction

2.1 A k -armed Bandit Problem

2.2 Action-value Methods

2.3 The 10-armed Testbed

2.4 Incremental Implementation

2.5 Tracking a Nonstationary Problem

2.6 Optimistic Initial Values

2.7 Upper-Condition-Bound Action Selection

2.8 Gradient Bandit Algorithms

2.9 Associative Search (Contextual Bandits)

2.0 Introduction

- The most important feature distinguishing reinforcement learning from other types of learning

uses training information that **evaluates the actions taken** rather than instructs by giving correct actions

o/n
reinforcement
→ actions evaluated *intentionally*
• Evaluative feedback vs instructive feedback
general domain feedback *general domain specific feedback*

- In this chapter we study the **evaluative aspect** of reinforcement learning in a simplified setting

K-armed bandit problem

2.1 A k-armed Bandit Problem

- **Problem definition**

1. a choice among k different options, or actions
2. After each choice you receive a numerical reward (chosen from stationary probability distribution)

- **Object**

To maximize the expected total reward over some time period

2.1 A k-armed Bandit Problem

- Each of the k actions has **an expected or mean reward** given that action is selected

$$q_*(a) \doteq \mathbb{E}[R_t \mid A_t = a].$$

- A_t : the action on time step t
 - R_t : the reward on time step t
 - $q_*(a)$: an expected or mean reward
 - $Q_t(a)$: the estimated value of action a at time step t
- Handwritten notes:*
- A bracket groups A_t and R_t with the note $\Rightarrow \theta / \mu \theta \quad \gamma / \mu \gamma \dots$
- A blue arrow points from the word "closer" to $Q_t(a)$, indicating it is closer to $q_*(a)$.

2.1 A k-armed Bandit Problem

- Greedy action vs ϵ -greedy

Greedy action whose estimated value is greatest

ϵ -greedy is non greedy action

82 step이 1기리다
-> Exploitation
-> Exploration

- Exploitation vs Exploration

Exploitation : to maximize the expected reward on the one step

Exploration : the greater total reward in the long run

-> balancing exploitation and exploration

2.2 Action-value Methods

- **True action-value** $q_*(a)$
the mean reward when the action is selected
How to estimate? Averaging the rewards “actually received”

- **Sample-average Method** $Q_t(a)$

$$Q_t(a) \doteq \frac{\text{sum of rewards when } a \text{ taken prior to } t}{\text{number of times } a \text{ taken prior to } t} = \frac{\sum_{i=1}^{t-1} R_i \cdot \mathbb{1}_{A_i=a}}{\sum_{i=1}^{t-1} \mathbb{1}_{A_i=a}},$$

where $\mathbb{1}_{\text{predicate}}$ denotes the random variable that is 1 if *predicate* is true and 0 if it is not.

If denominator is

- 1) $0 \rightarrow$ define $Q_t(a)$ as default value (such as 0)
- 2) $\infty \rightarrow Q_t(a)$ converges $q_*(a)$

2.2 Action-value Methods

- Greedy action selection

$$A_t \doteq \arg \max_a Q_t(a)$$

always exploits current knowledge to maximize immediate reward

- ϵ -greedy method

behave every once in a while with small probability ϵ

2.3 The 10-armed Testbed

- To compare the relative effectiveness of the greedy and ϵ -greedy
- A set of 2000 randomly generated k-armed bandit problems with $k = 10$ (normal distribution, mean 0, var 1)

2.3 The 10-armed Testbed




Figure 2.1: An example bandit problem from the 10-armed testbed. The true value $q_*(a)$ of each of the actions was selected according to a normal distribution with mean zero and unit variance, and then the rewards were selected according to a mean $q_*(a)$ unit variance normal distribution, as suggested by these distributions.

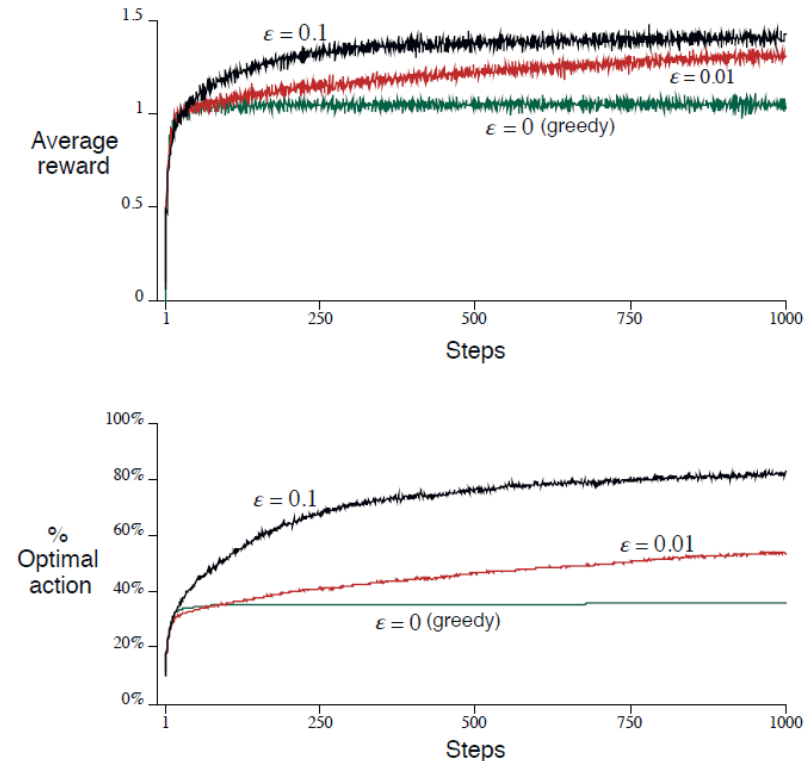


Figure 2.2: Average performance of ϵ -greedy action-value methods on the 10-armed testbed. These data are averages over 2000 runs with different bandit problems. All methods used sample averages as their action-value estimates.

2.4 Incremental Implementation

- Effective way that estimate action values as sample averages ?
- Q_n : estimate of its action value after it has been selected $n-1$ times

$$Q_n = \frac{R_1 + R_2 + \dots + R_{n-1}}{n - 1}.$$

- **Problem:**

- 1) memory : record of all the rewards
- 2) computation : perform computation whenever the estimated value was needed

-> Incremental Implementation

2.4 Incremental Implementation

$$Q_n \doteq \frac{R_1 + R_2 + \cdots + R_{n-1}}{n-1}.$$

$$Q_{n+1} = \frac{1}{n} \sum_{i=1}^n R_i \quad \text{NewEstimate} \leftarrow \text{OldEstimate} + \text{StepSize} \left[\text{Target} - \text{OldEstimate} \right].$$

$$= \frac{1}{n} \left(R_n + \sum_{i=1}^{n-1} R_i \right)$$

$$= \frac{1}{n} \left(R_n + (n-1) \frac{1}{n-1} \sum_{i=1}^{n-1} R_i \right)$$

$$= \frac{1}{n} (R_n + (n-1)Q_n)$$

$$= \frac{1}{n} (R_n + nQ_n - Q_n)$$

$$= Q_n + \frac{1}{n} [R_n - Q_n],$$



→ generalization

2.5 Tracking a Nonstationary Problem

- Stationary bandit problems - the reward probabilities do not change over time //
- Reinforcement learning problems that are effectively nonstationary //

Handwritten note: $\frac{1}{n} \sum_{t=1}^n R_t$

- **Non stationary Problem**

it makes sense to give more weight to recent rewards than to long-past rewards.

-> use a constant step-size parameter $\alpha \in (0,1]$

ex)

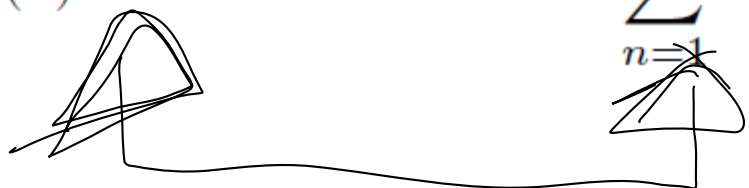
$$Q_{n+1} \doteq Q_n + \alpha [R_n - Q_n]$$

2.5 Tracking a Nonstationary Problem

- step-size parameter $\alpha \in (0,1]$

$$\begin{aligned} Q_{n+1} &= Q_n + \alpha [R_n - Q_n] \\ &= \alpha R_n + (1 - \alpha) Q_n \\ &= \alpha R_n + (1 - \alpha) [\alpha R_{n-1} + (1 - \alpha) Q_{n-1}] \\ &= \alpha R_n + (1 - \alpha) \alpha R_{n-1} + (1 - \alpha)^2 Q_{n-1} \\ &= \alpha R_n + (1 - \alpha) \alpha R_{n-1} + (1 - \alpha)^2 \alpha R_{n-2} + \\ &\quad \dots + (1 - \alpha)^{n-1} \alpha R_1 + (1 - \alpha)^n Q_1 \\ &= (1 - \alpha)^n Q_1 + \sum_{i=1}^n \alpha (1 - \alpha)^{n-i} R_i. \end{aligned}$$

- Convergence condition

$$\sum_{n=1}^{\infty} \alpha_n(a) = \infty \quad \text{and} \quad \sum_{n=1}^{\infty} \alpha_n^2(a) < \infty.$$


- Weighted average

$$(1 - \alpha)^n + \sum_{i=1}^n \alpha (1 - \alpha)^{n-i} = 1.$$

- Sample-average case

$$\alpha_n(a) = \frac{1}{n}$$

→ both

- Step-size parameter

$$\alpha \in (0,1]$$

→ second condition x

2.6 Optimistic Initial Values

1, 2, 3, 4, 5, 6, 7, 8, 9, 10

- $Q_1(a)$: initial action-value (initialize to 0)

if do not initialize $Q_1(a)$ to 0 → biased by initial estimate

→ used as a simple way to encourage exploration

“Optimistic Initial Values!”

ex) 10-armed bandit problem

Suppose initializing $Q_1(a)$ to +5 about every action a .

1. because of $Q_1(a) = 5$ about every action a , randomly select one action a
2. Whichever actions are initially selected, the reward is less than the starting estimates.
3. the learner switches to other actions (the selected action become non-greedy action)
4. repeat 1,2,3, about all actions

2.6 Optimistic Initial Values

- **Optimistic initial values**

effective on stationary problems, but not on non-stationary problem




Figure 2.3: The effect of optimistic initial action-value estimates on the 10-armed testbed. Both methods used a constant step-size parameter, $\alpha = 0.1$.

2.7 Upper-Condition-Bound Action Selection

- Exploration is needed because there is always uncertainty about the accuracy of the action-value estimates.
→ lets study exploration methods

- ϵ - greedy action selection forces the non-greedy actions to be tried, but no preference for those that are nearly greedy or particularly uncertain.

→ Upper-confidence-bound action selection(UCB)

Optimistic initial values

$$A_t \doteq \arg \max_a \left[Q_t(a) + c \sqrt{\frac{\ln t}{N_t(a)}} \right]$$

The increases get smaller over time, but are unbounded

→ All actions will eventually be selected

- $\ln t$: the natural logarithm of t
- $N_t(a)$: the number of times that action a has been selected prior to time t
- $c > 0$: the degree of exploration (confidence level, hyperparameter)
- UCB often performs well, but is more difficult than ϵ - greedy

2.7 Upper-Condition-Bound Action Selection

- UCB vs ϵ -greedy




Figure 2.4: Average performance of UCB action selection on the 10-armed testbed. As shown, UCB generally performs better than ϵ -greedy action selection, except in the first k steps, when it selects randomly among the as-yet-untried actions.

2.8 Gradient Bandit Algorithms

- **Method of learning a numerical preference (denote $H_t(a)$)**

Consider the relative preference of one action over another

- Action probabilities, which are determined according to a soft-max distribution

$$\Pr\{A_t = a\} \doteq \frac{e^{H_t(a)}}{\sum_{b=1}^k e^{H_t(b)}} \doteq \pi_t(a),$$

- $\pi_t(a)$: the probability of taking action a at time t
- Initially all preferences are the same ($H_1(a) = 0$)

$$\begin{aligned} H_{t+1}(A_t) &\doteq H_t(A_t) + \alpha(R_t - \bar{R}_t)(1 - \pi_t(A_t)), & \text{and} & & R_t - \bar{R}_t > 0 &\rightarrow \pi_t(a) & \uparrow \\ H_{t+1}(a) &\doteq H_t(a) - \alpha(R_t - \bar{R}_t)\pi_t(a), & & & R_t - \bar{R}_t < 0 &\rightarrow \pi_t(a) & \downarrow \end{aligned}$$

- Preferences are updated by the idea of stochastic gradient ascent.
- $\alpha > 0$: a step-size parameter
- \bar{R}_t : the average of all the rewards, as a baseline

2.8 Gradient Bandit Algorithms

- Baseline & step-size experiment

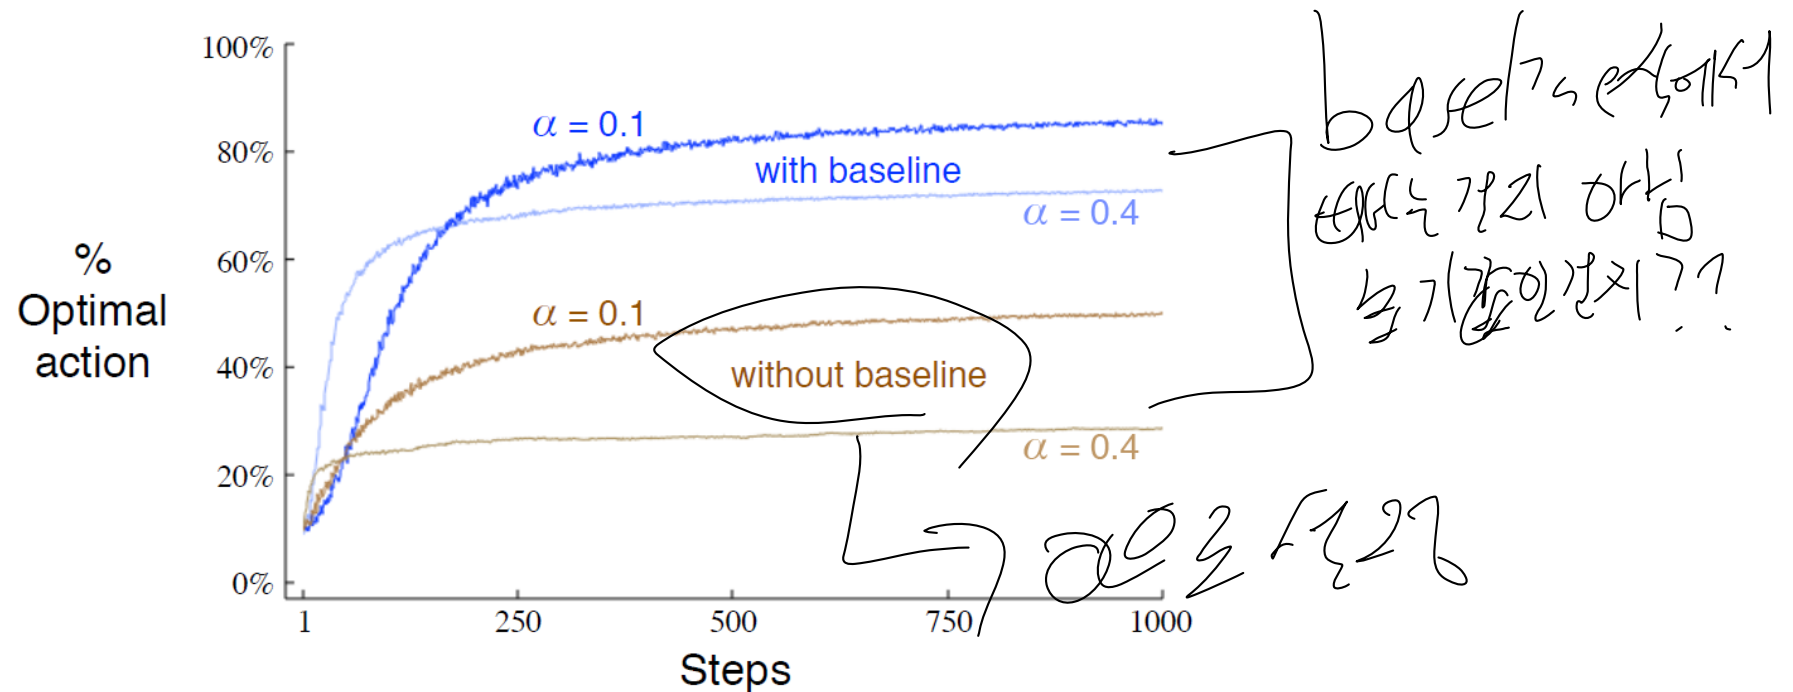


Figure 2.5: Average performance of the gradient bandit algorithm with and without a reward baseline on the 10-armed testbed when the $q_*(a)$ are chosen to be near +4 rather than near zero.

2.8 Gradient Bandit Algorithms

$$\begin{aligned}
 H_{t+1}(a) &\doteq H_t(a) + \alpha \frac{\partial \mathbb{E}[R_t]}{\partial H_t(a)}, \\
 \frac{\partial \mathbb{E}[R_t]}{\partial H_t(a)} &= \frac{\partial}{\partial H_t(a)} \left[\sum_b \pi_t(b) q_*(b) \right] \\
 &= \sum_b q_*(b) \frac{\partial \pi_t(b)}{\partial H_t(a)} \\
 &= \sum_b (q_*(b) - X_t) \frac{\partial \pi_t(b)}{\partial H_t(a)},
 \end{aligned}$$

$$\frac{\partial \mathbb{E}[R_t]}{\partial H_t(a)} = \sum_b \pi_t(b) (q_*(b) - X_t) \frac{\partial \pi_t(b)}{\partial H_t(a)} / \pi_t(b)$$

$$\begin{aligned}
 &= \mathbb{E} \left[(q_*(A_t) - X_t) \frac{\partial \pi_t(A_t)}{\partial H_t(a)} / \pi_t(A_t) \right] \\
 &= \mathbb{E} \left[(R_t - \bar{R}_t) \frac{\partial \pi_t(A_t)}{\partial H_t(a)} / \pi_t(A_t) \right], \\
 &= \mathbb{E} \left[(R_t - \bar{R}_t) \pi_t(A_t) (\mathbb{1}_{a=A_t} - \pi_t(a)) / \pi_t(A_t) \right] \\
 &= \mathbb{E} \left[(R_t - \bar{R}_t) (\mathbb{1}_{a=A_t} - \pi_t(a)) \right].
 \end{aligned}$$

$$H_{t+1}(a) = H_t(a) + \alpha (R_t - \bar{R}_t) (\mathbb{1}_{a=A_t} - \pi_t(a)), \quad \text{for all } a,$$

2.8 Gradient Bandit Algorithms

$$\begin{aligned}\frac{\partial \pi_t(b)}{\partial H_t(a)} &= \frac{\partial}{\partial H_t(a)} \pi_t(b) \\&= \frac{\partial}{\partial H_t(a)} \left[\frac{e^{H_t(b)}}{\sum_{c=1}^k e^{H_t(c)}} \right] \\&= \frac{\frac{\partial e^{H_t(b)}}{\partial H_t(a)} \sum_{c=1}^k e^{H_t(c)} - e^{H_t(b)} \frac{\partial \sum_{c=1}^k e^{H_t(c)}}{\partial H_t(a)}}{\left(\sum_{c=1}^k e^{H_t(c)} \right)^2} && \text{(by the quotient rule)} \\&= \frac{\mathbb{1}_{a=b} e^{H_t(b)} \sum_{c=1}^k e^{H_t(c)} - e^{H_t(b)} e^{H_t(a)}}{\left(\sum_{c=1}^k e^{H_t(c)} \right)^2} && \text{(because } \frac{\partial e^x}{\partial x} = e^x \text{)} \\&= \frac{\mathbb{1}_{a=b} e^{H_t(b)}}{\sum_{c=1}^k e^{H_t(c)}} - \frac{e^{H_t(b)} e^{H_t(a)}}{\left(\sum_{c=1}^k e^{H_t(c)} \right)^2} \\&= \mathbb{1}_{a=b} \pi_t(b) - \pi_t(b) \pi_t(a) \\&= \pi_t(b) (\mathbb{1}_{a=b} - \pi_t(a)).\end{aligned}$$

Q.E.D.

2.9 Associative Search (Contextual Bandits)

→ 0.52K6E2AX

- K-armed bandit problem → nonassociative task
- However, in a general reinforcement learning task there is more than one situation, and the goal is to learn a policy
- K-armed bandit problem → associative search task(contextual bandits) → full reinforcement learning.

→ (0.52K6E2AX 0.22/10/22)



THANK YOU